

Task 5.1, a

The learning approach described in the article uses deep learning in convolutional neural net for developing optimal $Q^*(s, a)$ action-value function for playing set of computer games in Atari 2600. Two key ideas were introduced for the learning: first one is using the pool of stored samples to take random and perform "replay" on them, second one is periodically updating target values, not always. For updating was used loss function

$$L_i(\theta_i) = E_{(s,a,r,s') \approx U(D)} [(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i))^2]$$

Here θ_i are the parameters (weights of the neural net) on the $i - th$ step, γ is a discount factor (that was taken 0.99) for learning, θ_i^- the parameters that are used for obtaining target values on the $i - th$ step (they are updated every C steps), $U(D)$ is a batch of experience drawn uniformly from stored experiences (tuples (s_t, a_t, r_t, s_{t+1})). The input for the convolutional neural net was $84 \times 84 \times 4$ image, preprocessed from the raw input from Atari 2600 and the output is a vector of $Q - values$ for each of the possible actions. Convolutional layers of Q- network here played role of feature extractor and were used for dimension reduction: instead of $4 \times 84 \times 84$ input nodes we got $32 \times 9 \times 9$ input nodes. For different games were set different sets of actions. Also, the rewards in the game were simplified to 1, 0 and -1, as the scaling was different in every game. For generating next action was used greedy policy, action is given to the framework Atari 2600 and next image and reward are returned. In general algorithm worked in the following sequence:

1. Make an action following greedy policy
2. Get the reward and next state, save them to the pool
3. Sample transitions from the pool and set the max value to calculate loss function
4. Perform gradient descent step
5. every C steps update weights for generating max Q values

For testing the results also the greedy policy with $\epsilon = 0.05$ was used in order to minimise the possibility of overfitting.

Task 5.1, b

States are the frames from Atari 2600, simplified slightly in order to reduce the memory and processing consumptions. So it is $84 \times 84 \times 4$ image of the current positions in the game. Action is a possible action in the specific game. They are set as initial values for the neural net, and represented by the output layer.

Task 5.1, c

Parameters that are adjusted by learning are the weights θ_i of the convolutional neural net (Q-network) at $i - th$ iteration. Number of the parameters (weights): $8 \times 8 \times 4 \times 16$

(1st hidden layer has 16 filters, and each contains $4 \times 8 \times 8$ weight matrixes - as we have four, not one frame as input) + $4 \times 4 \times 16 \times 32$ (2nd hidden layer: 32 filters, each has $16 \times 4 \times 4$ matrixes, for each input (filter from previous layer) + $9 \times 9 \times 32 \times 256$ (3rd hidden, that consist of 256 units, each has $32 \times 9 \times 9$ matrix - 32 - number of inputs from layer before) + 256×4 (output, 4 in case we have exactly 4 actions).

Task 5.2

[1]

References

- [1] *Deep Learning for Reinforcement Learning in Pacman*, <http://tuprints.ulb.tu-darmstadt.de/id/eprint/3832>, Bachelor-Thesis von Aaron Hochlißnder aus Wiesbaden Juli 2014.